

# AIOps

## Table of Contents

1. Part 1: Problem Statement
2. Part 2: Standard ML Solution
3. Part 3: Ops Challenges
4. Part 4: AIOps with mlflow
5. Part 5: Analysis and Interpretation

# Part I Problem Statement

## Problem Statement: Predicting Quality of Service (QoS) Metrics for 5G Network Optimization

### Overview

Telecom operators face challenges in maintaining optimal Quality of Service (QoS) for their users due to varying signal conditions, network congestion, and application-specific demands. The provided dataset includes time-series data on signal strength, latency, bandwidth requirements, allocated bandwidth, and resource allocation for different application types. The goal is to build a machine learning model to predict QoS metrics such as **latency** and **resource allocation efficiency**, enabling proactive network management and improved user experience.

---

### Objective

Develop a machine learning model to predict:

1. **Latency:** Estimate the latency for a given user, application type, and signal condition.
2. **Resource Allocation Efficiency:** Predict the ratio of allocated bandwidth to required bandwidth under varying network conditions.

The predicted values will help identify potential QoS degradation before it occurs, allowing for targeted network adjustments.

---

### Key Business Questions

1. **Latency Prediction:** Can we accurately predict latency based on signal strength, application type, and bandwidth requirements?
  2. **Resource Allocation Optimization:** Can we predict whether the allocated bandwidth will meet user demand efficiently?
  3. **Proactive QoS Management:** Can these predictions highlight potential QoS bottlenecks for specific users or applications?
-

## Detailed Problem Description

### 1. Latency Prediction:

- Latency is a critical metric for QoS, particularly for applications like emergency services and online gaming that require real-time responses.
- Variations in signal strength and bandwidth availability often cause unpredictable latency spikes.
- The model will predict latency based on:
  - Signal strength.
  - Application type.
  - Required and allocated bandwidth.

### 2. Resource Allocation Efficiency Prediction:

- Resource allocation efficiency measures how effectively network bandwidth meets application demands.
- An efficiency ratio below a threshold indicates potential resource under-provisioning, affecting QoS.
- The model will predict resource allocation efficiency using:
  - Application type.
  - Required bandwidth.
  - Allocated bandwidth.
  - Historical resource utilization patterns.

### 3. Data Patterns and Dependencies:

- Signal strength below a certain threshold may correlate with higher latency and lower resource allocation efficiency.
  - Specific application types (e.g., video calls) may consistently experience resource constraints.
  - Temporal patterns in data (e.g., peak vs. non-peak hours) could impact QoS metrics.
-

Dataset Fields and Their Roles

Field Name	Description	Role in Prediction
Timestamp	Date and time of the record	Temporal analysis (peak/non-peak times)
User_ID	Unique identifier for the user	Not directly used but useful for grouping if required
Application_Type	Type of application (e.g., video call, gaming)	Key feature for predicting QoS metrics
Signal_Strength	Signal strength in dBm	Predictor for latency and resource allocation efficiency
Latency	Observed latency in milliseconds	Target variable for latency prediction
Required_Bandwidth	Bandwidth required by the application	Predictor for resource allocation efficiency
Allocated_Bandwidth	Bandwidth allocated to the application	Predictor and part of target (efficiency calculation)
Resource_Allocation	Efficiency of resource allocation (0-1 scale)	Target variable for efficiency prediction

Potential Insights

- 1. **Latency Predictors:**
  - Identify key contributors to high latency for different application types (e.g., weak signal strength, high bandwidth demand).
- 2. **Bandwidth Efficiency Trends:**
  - Discover patterns of inefficient resource allocation and suggest improvements.
- 3. **Application-Specific QoS Challenges:**
  - Highlight applications that consistently experience QoS degradation and the factors contributing to it.

---

## **Assumptions and Limitations**

### **1. Assumptions:**

- Data is clean, consistent, and representative of the overall network usage patterns.
- Signal strength and bandwidth metrics are accurately recorded in real-time.

### **2. Limitations:**

- Geographic and user-specific metadata are not included, which could limit personalized predictions.
- Predictions will be based solely on the provided features without additional contextual data (e.g., network topology).

## Part 2: Proposed Solution

### Proposed Approach

#### 1. Exploratory Data Analysis (EDA):

- Identify trends, correlations, and anomalies in key features (e.g., signal strength, latency).
- Analyze the distribution of QoS metrics by application type and signal strength ranges.

#### 2. Feature Engineering:

- Create derived features:
  - **Bandwidth Efficiency:**  
$$\frac{\text{Allocated Bandwidth}}{\text{Required Bandwidth}}$$
  - **Signal Quality Category:** Group signal strength into ranges (e.g., strong, moderate, weak).
  - **Temporal Features:** Extract time-related features (e.g., hour of day, day of week).

#### 3. Machine Learning Models:

- **Supervised Learning:**
  - Target Variables:
    1. **Latency:** A regression task to predict latency in milliseconds.
    2. **Resource Allocation Efficiency:** A regression task to predict bandwidth efficiency as a ratio.
  - Input Features:
    - Signal strength, application type, required bandwidth, allocated bandwidth, and temporal features.
- **Candidate Models:**
  - Linear Regression, Random Forest, Gradient Boosting (e.g., XGBoost, LightGBM), and Neural Networks.
- Evaluate models using metrics like MAE (Mean Absolute Error) and RMSE (Root Mean Squared Error).

#### 4. Validation and Testing:

- Split data into training, validation, and test sets.
  - Use cross-validation to ensure robust performance metrics.
-

## Part 3: Current Challenges for Ops

### Instructions:

#### 1. Role-Playing

- Take on a role (e.g., Data Scientist, Operations Engineer, or Stakeholder).
  - Discuss challenges your role faces in ML workflows (e.g., tracking, deployment, monitoring). Relate to current use case.
  - Share key pain points and their impact.
- 

#### 2. Scenario Exploration

- Scenario: A deployed model's performance degrades after two weeks.
  - Discuss:
    - Possible causes of degradation.
    - Metrics needed to investigate.
    - Solutions to prevent this in the future.
- 

#### 3. Pain Point Mapping

- List the steps in the ML workflow (e.g., data prep, training, deployment).
  - Identify pain points for each step (e.g., manual processes, lack of tracking).
  - Share how addressing these could improve outcomes.
- 

#### 4. Challenge Prioritization

- Review and rank the top three challenges from discussions.
- Explain why they're critical and suggest possible solutions.
- Discuss tools or practices (e.g., tracking systems, automation) to resolve them.

**1. Relate Challenges to Your Experience:**

- Reflect on any past machine learning projects you've worked on. Think about the pain points during model development, deployment, or monitoring.
- Discuss examples of manual processes or inefficiencies that could benefit from automation or tracking.

**2. Break Down the Workflow:**

- Think about the end-to-end machine learning lifecycle, from data preparation to deployment. Identify areas where tracking, monitoring, or scaling were difficult.

**3. Focus on Outcomes and Impact:**

- Consider what happens when models fail in production.
- Discuss scenarios where issues like lack of reproducibility, inconsistent results, or system bottlenecks negatively impacted project outcomes.

**4. Challenge Current Practices:**

- Explore if the current methods provide enough transparency, scalability, and collaboration.
- Think critically about what information was missing when troubleshooting or improving a model.

**5. Prioritize Scalability and Maintenance:**

- Imagine your solution being deployed in a real-world, high-scale environment. Discuss potential risks or challenges in maintaining the system.



## Hint (Examples):

### Experimentation and Tracking

1. *"How do you currently keep track of the parameters, datasets, and metrics used during training? Have you ever struggled to reproduce a result from a previous experiment?"*
2. *"What happens when someone on the team modifies a model or pipeline? How do you ensure these changes are tracked and logged?"*

### Model Versioning and Governance

3. *"Imagine deploying an updated model to production and it performs worse than the previous version. How would you roll back to the earlier version? Do you currently have a system for managing multiple versions?"*
4. *"How do you ensure that every version of a model is linked to the exact dataset and code used to train it?"*

### Monitoring Operational Metrics

5. *"Once your model is deployed, what operational metrics would you track to ensure smooth functioning? Would you monitor things like CPU usage, memory, or latency?"*
6. *"How do you know if your deployed model is drifting in terms of accuracy or encountering degraded performance over time?"*

### Collaboration and Transparency

7. *"How do team members currently share experiment results, model artifacts, or deployment processes? Is it easy to collaborate across roles?"*
8. *"Have you ever faced issues where one team member's changes to a pipeline broke something for someone else? How could better transparency have helped?"*

### Automation and Workflow Management

9. *"What aspects of your current machine learning workflow are repetitive or manual? How could automating these processes help improve efficiency?"*
10. *"Imagine retraining and deploying a model every week with new data. What would make this process more streamlined and reliable?"*

### Scaling and Real-Time Use

11. *"What happens when your model needs to handle 10x more data or user queries? Have you thought about how resource constraints might affect scalability?"*
12. *"If you needed to deploy multiple models for different tasks (e.g., latency and efficiency), how would you manage dependencies and scaling?"*

### Success Metrics

13. *"How do you currently measure the success of your deployed models? Are these metrics tracked consistently across experiments and in production?"*

14. *"Have you ever struggled to define clear performance metrics for a model? How could a standardized system help?"*